
Understanding Grammar Representation in the Human Brain via Information-Restricted LLMs

Bachelor's Thesis

Author

Simone Esposito
Matrikelnummer: 469561

Supervisors

Prof. Dr. Fatma Deniz
Prof. Dr. Klaus Obermayer

Berlin, February 2026

Abstract

This thesis investigates the brain representation of grammatical class (GC) in the human brain, focusing on nouns and verbs. Using voxelwise encoding models (VEM) and contextual embeddings from GPT-2, I examine whether the brain processes nouns and verbs in distinct areas. By applying attention masks to generate noun-masked and verb-masked embeddings, I fit these features to fMRI data collected from subjects listening to and reading naturalistic stories. The results reveal a complex pattern of GC representation that challenges traditional dichotomies and supports a more distributed account of language processing in the brain.

Zusammenfassung

Diese Arbeit untersucht die neuronale Repräsentation von grammatischer Klasse (GC) im menschlichen Gehirn, mit Fokus auf Nomen und Verben. Mithilfe von voxelweisen Enkodierungsmodellen (VEM) und kontextuellen Einbettungen von GPT-2 analysiere ich, ob das Gehirn Nomen und Verben in unterschiedlichen Bereichen verarbeitet. Durch die Anwendung von Aufmerksamkeitsmasken zur Generierung von Nomen- und Verb-maskierten Einbettungen passe ich diese Merkmale an fMRI-Daten an, die von Probanden gesammelt wurden, während sie natürliche Geschichten hörten und lasen. Die Ergebnisse zeigen ein komplexes Muster der GC-Repräsentation, das traditionelle Dichotomien herausfordert und eine stärker verteilte Darstellung der Sprachverarbeitung im Gehirn unterstützt.

Introduction

Language is a uniquely human capacity, and grammar is its organizing principle. Among the most fundamental categories of grammar are parts of speech (POS). While these include various word classes such as adjectives and adverbs, the distinction between nouns and verbs has remained the primary focus of research in psycholinguistics and cognitive neuroscience. To investigate the functional organization of this distinction, researchers have historically utilized neuropsychological case studies of lesion-deficit patterns and, more recently, brain imaging techniques such as functional magnetic resonance imaging (fMRI)."

Neuropsychological studies

Early neuropsychological studies conducted experiments with aphasic patients who had strong impairments in the retrieval and production of nouns and verbs respectively. A. R. Damasio and Tranel (1993), for example, observed impairments in noun retrieval for two patients with damage to the *left anterior and middle temporal lobe*, while the *left premotor cortex* was named as a possible cause of the third patient's impairment in the retrieval of verbs. This observation gave rise to the *fronto-temporal* dichotomy, suggesting that nouns might be processed in temporal areas while verbs are processed in frontal areas. According to that study, the distinction between nouns and verbs can be attributed to the semantic properties of the two classes, with nouns being more associated with objects and verbs with actions.

However, subsequent studies have found contradicting results, with some patients showing frontal damage without verb deficits and temporal damage without noun deficits (H. Damasio et al. 2001; De Renzi and Pellegrino 1995). This led to the examination of other factors that might contribute to the observed dissociation, such as morphological inflection, which is more common in verbs than nouns (K. Shapiro and Caramazza 2003; Tyler et al. 2004) or imageability (Bird et al. 2000).

Brain Imaging Studies

With the advancement of brain imaging techniques, such as fMRI, researchers have started to investigate the cortical correlates of noun and verb processing in healthy individuals. The results of these studies can be broadly categorized into three groups. First, some studies have found the distinction between nouns and verbs in the brain to be mainly driven by the semantic properties of the words, with nouns being more associated with objects and verbs with actions (Vigliocco et al. 2011). This view is often associated with the **emergentist** perspective in psycholinguistics, which posits that noun and verb distinctions emerge from semantic constraints rather than being explicitly represented in the brain (Vigliocco et al. 2011).

Second, some studies have suggested that the difference can be attributed to grammatical operations like morphological inflection, which are more common in verbs than nouns. This view is often associated with the **combinatorial** perspective, which posits that noun and verb distinctions are procedural and

related to the combinatorial properties of words (Siri et al. 2007; Yokoyama et al. 2006; Tyler et al. 2004).

Finally, some meta-analytical evidence suggests a lack of consistent spatial segregation between nouns and verbs across different experimental tasks (Crepaldi et al. 2013), pointing to a more **distributed** representation of nouns and verbs in the brain (Crepaldi et al. 2013).

The main limitation of these studies is that they rely on highly controlled stimuli, such as single words or simple sentences, which may not capture the complexity of natural language processing.

cite and tell which study uses what. this is too general.

Naturalistic Studies

To overcome the limitations of controlled stimuli, recent research has shifted toward naturalistic studies. These utilize ecological stimuli such as continuous stories or spontaneous speech, which better reflect real-world language processing (A. G. Huth et al. 2016). To analyze the complex brain responses elicited by such stimuli, Voxelwise Encoding Models (VEM) have emerged as a primary methodological framework. VEMs allow researchers to predict the BOLD response of individual voxels in the brain based on features extracted from naturalistic stimuli, such as stories or conversations. This approach has been used to investigate various aspects of language processing, including syntax, semantics, and stimulus modality (A. G. Huth et al. 2016; Deniz et al. 2019). Particularly with the advancement of large language models (LLMs) such as GPT-2 and their ability to generate contextualized embeddings, VEMs have become a powerful tool to investigate the cortical representation of language in the brain.

I used VEMs to investigate the representation of nouns and verbs in the brain. After masking the attention of GPT-2 to either nouns or verbs during the generation of contextualized embeddings, I was able to create information-restricted feature-spaces. I then fit these feature spaces to the BOLD responses of 9 subjects who listened to and read naturalistic stories. Finally, I compared the performance of the noun-masked and the verb-masked model across the brain, to identify areas that were more sensitive to nouns and verbs respectively.

there is not enough content or intro given yet to be able to understand this.

The central question of this thesis is whether the brain represents nouns and verbs in distinct areas, and if so, which of the three theoretical accounts — emergentist, combinatorial, or distributed — best characterises the observed pattern.

Materials and Methods

Stimuli

The speech stimuli consist of 11 short stories taken from **The Moth Radio Hour**, a storytelling program where true autobiographical stories are told in front of a live audience. These were used in previous papers (Huth et al., 2016), including the one where the response data was gathered (Deniz et al., 2019) providing a naturalistic stimulus.

Out of these 11 stories, 10 were used as training data while one 10-minute

mention that this is fmri data, taken from this study. modalities, number of people....

Table 1: Theoretical accounts of grammatical class representation: predicted patterns versus observed results.

Theoretical account	Predicted pattern	Observed pattern
Fronto-Temporal View <i>Neuropsychological</i> (A. R. Damasio and Tranel 1993)	V>N in frontal regions; N>V in temporal regions	V>N in temporal regions (ATL, PTL); no clean frontal/temporal split.
Emergentist View <i>Object-Action, Imageability</i> (Bird et al. 2000; Vigliocco et al. 2011)	V > N in frontal and parietal regions; N > V in inferior temporal regions.	N > V near visual areas; Similar sensitivity in frontal regions.
Combinatorial <i>Morphology/inflection</i> (Tyler et al. 2004; K. A. Shapiro et al. 2005)	N = V broadly, with V > N specifically in IFG for morphological composition.	Slight V > N in IFG; N > V in IFGOrb.
Distributed <i>Shared network</i> (Crepaldi et al. 2013; Caucheteux et al. 2021)	Noun and verb circuits distributed across frontal, parietal, and temporal regions.	Smaller clusters of noun and verb sensitive voxels across the brain.

story was used as validation data. The latter was played two times during the experiment (in two separate runs) to which the responses were finally averaged across the two runs.

fix sentence structure

fMRI Data

The fMRI data used in this thesis were originally collected and preprocessed by Deniz et al. (2019). The dataset consists of brain responses from 9 subjects who both listened to and read the stimuli in separate sessions. Each spoken (and written) story was presented during a separate fMRI scan. The sampling rate was 2 seconds, so 1 Repetition Time (TR) = 2 s. Each scan included 10 seconds (5 TR) of silence both before and after the story. These data were collected during two 3-hour scan sessions that were performed on different days. MRI data were collected on a 3T Siemens TIM Trio scanner at the UC Berkeley Brain Imaging Center.

Each functional run was motion-corrected using the FMRIB Linear Image Registration Tool (FLIRT) from FSL 5.0 (Jenkinson and Smith, 2001; Jenkinson et al., 2002). All volumes in the run were then averaged over time to obtain a high-quality template volume. FLIRT was also used to automatically align the template volume for each run with the overall template, which was chosen to be the temporal average of the first functional run for each participant.

Features

Semantic features from English1000 and several low level features from Deniz et al. (2019) were used. Additionally, embeddings from an LLM were created.

English1000

The English1000 feature space is a 985-dimensional semantic model based on word co-occurrence statistics (Deerwester et al. 1990; Lund and Burgess 1996; Mitchell et al. 2008; A. G. Huth et al. 2016). A co-occurrence matrix M was constructed with 985 rows corresponding to basic English words (drawn from Wikipedia’s List of 1000 Basic Words) and the columns corresponding to the words selected from a large text corpus comprising Moth Radio Hour transcripts, books from Project Gutenberg, Wikipedia articles, and Reddit comments (A. G. Huth et al. 2016).

Each entry M_{ij} in the matrix represents the number of times the i -th basic word co-occurs with the j -th word in the corpus within a 15-word window. Counts were then log-transformed as $\log(1 + M_{ij})$, and rows and columns were each z-scored to correct for differences in word frequency.

The feature vector for a given word w is then obtained by looking up the corresponding column in the matrix, yielding a 985-dimensional vector that captures the semantic relationships of w with the basic English words. These features were used when fitting the model to the listening data. For more details on the construction of the English1000 feature space, see Deniz et al. (2019) and A. G. Huth et al. (2016).

Phonemes

A list of phoneme-time pairs was constructed for each of the 39 English phonemes. This list was then down-sampled to the fMRI acquisition rate. The result is a $(39 \times N_{TR})$ array. These features were used when fitting the model to the listening data.

Letters

Similar to the phonemes, a 26-parameter model was created to represent the frequency of each of the 26 letters in the English alphabet that appeared in the written stories. This was also down-sampled to the sampling rate, resulting in a $(26 \times N_{TR})$ array.

Additional Features

To account for the highly variable speech rate both within and across stories, single-feature models were constructed. These count the number of words, number of phonemes, number of letters, and number of story speaker’s pauses that occurred during the acquisition of each fMRI volume (2.0045 s). To account for the variable word lengths during the visual presentation, a single-feature word length variation model was constructed by taking the variance of word lengths that occurred during the acquisition of each fMRI volume.

mention
all low
level to-
gether, tell
which are
specific
to which
modality.

Contextual Embeddings

The largest issue with the lexical *English1000* embeddings is the lack of contextual information in the feature space. Put differently, the same word will be mapped to the same feature vector regardless of the context it appears in. For example, the word *fine* has a very different meaning in the sentences “I am fine” and “I paid a fine”.

To account for this, I made use of *contextual embeddings*, which, as the name suggests, aim to include contextual information in the word embeddings. Such embeddings were introduced by the Long Short-Term Memory (LSTM) network ELMo (Peters et al. 2018) and refined by the transformer model BERT (Devlin et al. 2019).

Contextual embeddings play a fundamental role in all modern Language Models (LM), and for this thesis I decided to use GPT-2’s embeddings, as it is commonly used in related literature (Schrumpf et al. 2021; Tuckute et al. 2024).

Given an input sequence comprised of the i -th word of a story w_i and its k -length context $C_i = (w_s, \dots, w_{i-1})$ with $s = \max(0, i - k)$, I can generate the contextual embeddings through the following steps.

Tokenization

Before a sentence or word can be embedded into vectors, it needs to be broken up into a set of token IDs. This turns the input sequence into a sequence of tokens $(t_1, \dots, t_L) \in \mathbb{N}^L$.

Initial Embeddings

Upon receiving the token sequence, GPT-2 first converts the tokens into embedding vectors via a *lookup table* with a fixed embedding size $d_{\text{model}} = 768$, then adds positional encodings to generate the initial embeddings. These initial embeddings are non-contextual.

Hidden Layers

The initial embeddings then go through a sequence of layers where they are first weighted using the attention mechanism (Vaswani et al. 2023), then passed through a feed-forward network to enrich them with contextual information from surrounding tokens. This yields a new contextualized embedding of the same dimensionality d_{model} .

Previous studies have found that the middle layers, e.g., Layer 8 for GPT-2 small, are best at predicting brain activity in VEMs (Schrumpf et al. 2021; Jain and A. Huth 2018), and Layer 8 is therefore used throughout this work.

Embedding Pooling

The hidden state $H_8 \in \mathbb{R}^{L \times 768}$ contains the contextualized embeddings for every token in the input sequence. These tokens are then averaged to return w ’s contextualized embedding e_w :

$$e_w = \frac{1}{L} \sum_{i=1}^L h_8^{(i)}$$

Downsampling

After repeating this process for each word in each story, I have an embedding matrix $E_j \in \mathbb{R}^{N_j \times 768}$ with $N_j = |S_j|$ being the number of words in the j -th story. To fit the features to the fMRI data, the contextual embeddings are downsampled onto a feature space $F_j \in \mathbb{R}^{K_j \times 768}$ via a Lanczos filter set to the




Figure 1. Contextual embedding pipeline with POS attention masking. A stimulus sentence is passed through a noun filter, which uses spaCy POS tags to generate an attention mask (green: unmasked; red: masked). The masked input is fed into GPT-2, whose hidden state at Layer 8, $H_8 \in \mathbb{R}^{L \times 768}$, is then averaged over unmasked tokens to produce a single contextualized embedding $e_w \in \mathbb{R}^{768}$ for the target word.

fMRI Nyquist rate (0.25 Hz) with K_j being the number of sample points (TRs) of the j -th story.

The baseline contextualized embeddings are generated using a context length of 20 words for each word in each story.

Information Restricted LLMs

The masking of Parts of Speech (POS) was performed automatically using Python’s `spaCy` library. This library first splits an input sequence into a series of word-like tokens based on language-specific rules. For example, *don’t* is split into two tokens *do* and *n’t*. These tokens are then assigned a POS tag that takes their context into consideration.

Given a set of POS tags P , I can generate a boolean mask for each story. (Fig 1)

$$M_j = (f_{\text{mask}}(t_{j,1}), \dots, f_{\text{mask}}(t_{j,n})) \quad f_{\text{mask}}(t) = \begin{cases} 0 & \text{POS}(t) \in P \\ 1 & \text{POS}(t) \notin P \end{cases}$$

Since proper nouns and auxiliary verbs perform similar functions to nouns and verbs respectively, they are included in the noun and verb POS categories accordingly.

Attention Mask

To apply the POS masks during the generation of contextual embeddings for a word w , the mask values for w and its context C are first mapped to the GPT-2 tokens, yielding M'_w . They are then used as *attention masks*, indicating which tokens the model should attend to and which should not. The resulting contextual embedding is calculated by averaging only the unmasked hidden states:

$$\bar{e}_w = \frac{1}{L'} \sum_{i=1}^L m_w'^{(i)} h_8^{(i)} \quad L' = \sum_{i=1}^L m_w'^{(i)}$$

Rather than removing masked tokens from the input sequence — which would alter positional encodings and shift the attention context of remaining tokens — attention masking preserves the original sequence structure while selectively excluding the target POS category from the model’s attention computation.

Voxelwise Encoding Model Framework

Voxelwise encoding models (VEM) aim to predict the BOLD response for each individual voxel v from the stimulus S . This is usually done by first extracting one or more sets of features (called feature spaces) from the stimulus, then modeling the brain response as a linear combination of those feature spaces.

This corresponds to approximating the function $f(S) = g(S)\beta$ where g transforms the N -length stimulus into an $N \times P$ feature space with β as the $P \times V$ weights, V being the number of voxels. This can be accomplished by fitting the feature spaces $g(S_{\text{train}})$ to the responses R_{train} of a training dataset, yielding a model $\hat{f}$. The hemodynamic delay is accounted for by including delays of 0, 4, 8, and 16 seconds for each feature, resulting in a finite impulse response (FIR) model.

Finally, the model accuracy is calculated as the per-voxel Pearson’s correlation r_v between the model predictions on a validation dataset $\hat{R}_{\text{test}} = \hat{f}(S)$ and the actual responses R_{test} .

The value of r_v can then be interpreted as the extent to which voxel v ’s response to the stimulus can be explained by the given feature space. Projecting r_v onto the cortical surface then yields a sensitivity map to the specific feature space. This shifts the focus from experimental design to the choice of features, allowing researchers to test different hypotheses utilizing the same dataset by varying the feature space.

Banded Ridge Regression

To fit the multiple feature spaces to the BOLD responses I used banded ridge regression (BRR), which is based on Ridge Regression, a linear regression method that augments the least-squares objective with an ℓ_2 penalty to reduce overfitting to the training data. This corresponds to finding the weight vector

$$\hat{\beta} = \arg \min_{\beta} \|X\beta - y\|_2^2 + \alpha \|\beta\|_2^2$$




Figure 2. Parts of Speech sensitivity across the brain using all POS: nouns, verbs, adjectives, adverbs, pronouns and interjections. **a** Voxelwise model r values in the listening modality using POS-masked contextual embeddings were subtracted from the r values of the baseline contextual model to get the drop in accuracy. The drop in accuracy for each POS was then ranked for each voxel individually and the greatest drop was plotted on the flatmap in the respective color **b** Voxels are colored using the same procedure as in **a** but for the reading modality. **c** Distribution of POS tags in the training stories and the validation story. The percentage of each POS tag was calculated by counting the number of words with that tag and dividing it by the total number of words in the story. The average percentage across the 10 training stories is plotted on the left, while the percentage for the validation story is plotted on the right.

with $\alpha > 0$ being the regularization parameter, which is typically selected through grid search or random search with cross-validation.

BRR extends Ridge Regression by assigning a separate regularization parameter α_i to each feature space X_i (Nunez-Elizalde et al. 2019).

$$\hat{\beta} = \arg \min_{\beta} \left\| \sum_{i=1}^m X_i \beta_i - y \right\|_2^2 + \sum_{i=1}^m \alpha_i \|\beta_i\|_2^2$$

This is important since the feature spaces can vary greatly in dimensionality and predictive power and thus need different levels of regularization.

Results

The goal of this thesis was to determine whether the brain represents nouns and verbs in distinct areas, and which theoretical account best characterises the observed pattern. To answer this question, I generated noun-masked and verb-masked contextual embeddings from 11 stories taken from *The Moth Radio Hour*. These embeddings were obtained by applying attention masks to the GPT-2 input during contextualized embedding generation, selectively suppressing either noun or verb tokens. I then fit these features, in addition to acoustic, phonetic and [what] features to the BOLD responses of 9 subjects who listened to and read the stories, using banded ridge regression. Finally, I compared the performance of the noun-masked and verb-masked model across the brain, to identify areas that were more sensitive to nouns and verbs respectively.

Contextual embeddings predict brain activity across the cortex

To ensure that the contextual embeddings were successfully able to predict brain activity, I first fit the unmasked contextual embeddings to the responses. After fitting the model, I calculated the per-voxel Pearson’s correlation r_v between the model predictions on the validation dataset and the actual responses. The resulting r_v values were then projected onto the cortical surface to yield a sensitivity map to the contextual embeddings. The contextual embeddings were able to predict brain activity in a large portion of the cortex. The average r_v across the cortex was 0.21 for the listening modality and 0.20 for the reading modality, with some voxels reaching values as high as 0.65. The standard error of the mean (SEM) across subjects was 0.0028 for the listening modality and 0.0021 for the reading modality, indicating that the results were consistent across subjects. This confirms the validity of the contextual embeddings as a feature space for predicting brain activity in response to naturalistic language stimuli.

POS-masked contextual embeddings reveal clusters of POS sensitivity across the brain

Before focusing on the noun-masked and verb-masked models, I first looked at the performance of a multitude of POS-masked features, including adjectives,

adverbs, pronouns and interjections. For each POS-masked model, I calculated the drop in performance compared to the baseline unmasked model by subtracting the r_v values of the masked model from the r_v values of the baseline model. I then ranked the drop in performance for each POS category for each voxel individually and plotted the greatest drop on the flatmap in the respective color. This allowed me to identify which areas of the brain were more sensitive to each POS category (Fig 2a,b).

The results showed clear clusters of POS sensitivity across the brain with nouns being more prevalent in the Medial Parietal Cortex (MPC), while verbs were consistently present in the Temporo-Parietal Junction (TPJ). Nouns and verbs resulted in the greatest drop in performance, followed by pronouns and adverbs, while adjectives and interjections had the least impact on the encoding performance.

The distribution of POS tags in the training stories and the validation story was also analyzed, to verify that the validation story is a representative sample of the training data and that any differences in model accuracy between POS categories reflect the brain’s sensitivity to grammatical class rather than an artifact of unequal stimulus composition. (Fig 2c).

Nouns and verbs show distinct patterns of sensitivity across the brain

To investigate the sensitivity to nouns and verbs in more detail, I focused on the noun-masked and verb-masked models. For each voxel, I calculated the drop in performance compared to the baseline unmasked model for both the noun-masked and verb-masked models. I then plotted the respective drops in performance for each voxel on the flatmap using a bivariate color map, where the x-axis (orange) represents the drop in performance for the noun-masked model and the y-axis (blue) represents the drop in performance for the verb-masked model. This allowed me to identify areas of the brain that were more sensitive to nouns (points above the diagonal) and areas that were more sensitive to verbs (points below the diagonal), as well as areas that were sensitive to both (points on the diagonal) (Fig 3). The results showed a clear separation between noun-sensitive and verb-sensitive areas. Moreover, the greatest drops for both nouns and verbs were observed in high-level processing areas, such as the MPC, TPJ and frontal areas.

Noun sensitive voxels seem to be especially concentrated around the Medial and Inferior Parietal Cortex (MCP and IPC), while verb sensitive voxels are more prevalent in the Temporo-Parietal Junction (TPJ). In the frontal region, the clusters appear to be smaller and more interleaved. The distribution of noun-sensitive and verb-sensitive voxels was consistent across subjects and modality, with some variability in the exact location and extent of the clusters.

Verb sensitivity is more prevalent across language ROIs

To further quantify the sensitivity to nouns and verbs across different brain regions, I used a set of Regions of Interest (ROIs) based on regions that are known to be involved in language processing (Negi et al. 2025; Fedorenko et al.




Figure 3. Noun and verb sensitivity across subjects in the listening modality. Voxelwise model r values in the listening modality for noun-masked and verb-masked contextual embeddings were subtracted from the r values of the baseline contextual model to get the drop in accuracy, then plotted using a bivariate color map. Blue points indicate a greater drop in performance for the noun-masked model while orange points indicate a greater drop for the verb-masked model. Grayscale points indicate a similar drop for both models. The large panel **a** shows the results for subject 07, while the smaller panels **b** show the results for the remaining subjects.




Figure 4. Noun and verb sensitivity across language ROIs. Each point indicates the mean drop in predictive performance $\Delta r = r_{\text{baseline}} - r_{\text{masked}}$ for the noun-masked (blue) and verb-masked (orange) models, averaged across subjects within each ROI. For each subject, Δr was computed as the mean positive drop across voxels within the ROI that survived FDR correction. Error bars represent ± 1 SEM across subjects. ROIs are ordered by their mean Δr in the listening modality.

2010). I first mapped each subject to FreeSurfer’s fsaverage template (Fischl et al. 1999). Then, I used the following ROIs: Middle Frontal Gyrus (MFG), Posterior Temporal Lobe (PTL), Angular Gyrus (AG), Anterior Temporal Lobe (ATL), Inferior Frontal Gyrus (IFG), Dorsomedial Prefrontal Cortex (dmPFC), Posterior Cingulate Cortex (PCC), and Inferior Frontal Gyrus, Orbital portion (IFGOrb). For each ROI, I then calculated the drop in performance for the noun-masked and verb-masked models $\Delta r = r_{\text{baseline}} - r_{\text{masked}}$ ($\Delta r \geq 0$) across voxels within the ROI where the drop was positive and calculated the average drop for each subject. Finally, I calculated the mean and standard error of the mean across subjects.

The results showed a consistent tendency toward greater verb sensitivity across ROIs, with the mean drop in performance ranging from 0.03 to 0.06 across both nouns and verbs for both listening and reading. The MFG and PTL showed a consistently larger drop for the verb-masked model across both modalities, suggesting that these regions rely more heavily on verb-specific contextual information. A greater drop in performance for the verb-masked model was also observed in the AG in the reading modality and in the ATL for the listening modality (Fig 4).

Discussion

In this thesis, I aimed to investigate the representation of Parts of Speech (POS) in the human brain, focussing on nouns and verbs. Using voxelwise encoding models (VEM) and POS-masked contextual embeddings from GPT-2, I was able to create a sensitivity map for nouns and verbs across the brain. The results revealed a pattern of segregated but distributed representations of nouns and verbs that challenges traditional dichotomies and supports a more distributed account of language processing in the brain.

The observed pattern of noun and verb sensitivity across the brain can be evaluated against the three theoretical accounts: the classic fronto-temporal account, the semantic account, the combinatorial account, and the distributed account.

The fronto-temporal dichotomy stems primarily from neuropsychological studies of aphasic patients and suggests that nouns are processed in temporal areas while verbs are processed in frontal areas. However, the results of this thesis show a higher sensitivity to verbs in temporal areas such as the ATL and AG, which is in direct contradiction to the fronto-temporal account. Moreover, frontal areas such as the IFG, dmPFC and IFGOrbs do not show a clear preference for verbs over nouns, instead showing a mixed pattern of sensitivity. Therefore, the results of this thesis do not support the classic fronto-temporal account of noun and verb representation.

This discrepancy can be mainly attributed to the difference in methodology. Neuropsychological studies rely on patients with brain damage, which usually extends over large areas and can affect multiple cognitive functions, making it difficult to isolate the specific contribution of a given area to noun and verb processing (Crepaldi et al. 2013). In contrast, VEMs allow for a more fine-grained analysis of the cortical representation of language by fitting specific

features to the BOLD response of individual voxels, thus providing a more direct measure of the sensitivity to nouns and verbs across the brain.

The semantic account posits that the distinction between nouns and verbs is mainly driven by the semantic properties of the words, with nouns being more associated with objects and verbs with actions.

The combinatorial account posits that the difference can be attributed to grammatical operations like morphological inflection, which are more common in verbs than nouns. Indeed, a slightly greater verb sensitivity observed in the IFG (Fig 4) could be consistent with this account, as the IFG has been implicated in morphological processing (Siri et al. 2007; Yokoyama et al. 2006; Tyler et al. 2004). This account, however fails to explain the greater verb sensitivity observed in the ATL and AG, which are not typically associated with morphological processing. Moreover, my methodology does not specifically test for morphological processing, as the contextual embeddings capture a wide range of linguistic information beyond morphology, including semantics and syntax.

The distributed account posits that both nouns and verbs are broadly represented across frontal, parietal and temporal areas (Crepaldi et al. 2013). The results of this thesis are consistent with this account, as both noun-sensitive and verb-sensitive areas were observed across the brain. Particularly in the frontal region, the clusters of noun and verb sensitivity appear to be smaller and more interleaved, which is consistent with a distributed representation. However, the strong separation between noun-sensitive and verb-sensitive areas observed in the parietal and temporal regions challenges a purely distributed account and suggests that there might be some degree of functional specialization for nouns and verbs in these areas. Therefore, the results of this thesis support a more nuanced version of the distributed account, where both nouns and verbs are broadly represented across the brain but with some degree of segregation in certain regions.

The contradictions between the results of this thesis and brain imaging studies can again be attributed to the difference in methodology. The use of naturalistic stimuli and VEMs allows for a more ecologically valid investigation of language processing in the brain, which may reveal patterns that are not apparent in highly controlled experimental settings like with block designs or task-based paradigms. In these controlled settings, participants are often required to perform specific tasks that may not reflect natural language processing, and the stimuli are usually limited to single words or simple sentences.

In conclusion, this thesis provides evidence for a complex representation of nouns and verbs in the human brain, with some degree of segregation in certain regions. This challenges traditional dichotomies and supports a more nuanced view of language processing in the brain that takes into account both distributed representations and functional specialization.

Limitations and Future Directions

References

- Bird, Helen, David Howard, and Sue Franklin (May 1, 2000). “Why Is a Verb Like an Inanimate Object? Grammatical Category and Semantic Category Deficits”. In: *Brain and Language* 72.3, pp. 246–309. ISSN: 0093-934X. DOI: 10.1006/brln.2000.2292. URL: <https://www.sciencedirect.com/science/article/pii/S0093934X00922922> (visited on 04/16/2026).
- Caucheteux, Charlotte, Alexandre Gramfort, and Jean-Remi King (July 1, 2021). “Disentangling syntax and semantics in the brain with deep networks”. In: *Proceedings of the 38th International Conference on Machine Learning*. International Conference on Machine Learning. PMLR, pp. 1336–1348. URL: <https://proceedings.mlr.press/v139/caucheteux21a.html> (visited on 02/19/2026).
- Crepaldi, Davide et al. (June 27, 2013). “Clustering the lexicon in the brain: a meta-analysis of the neurofunctional evidence on noun and verb processing”. In: *Frontiers in Human Neuroscience* 7. ISSN: 1662-5161. DOI: 10.3389/fnhum.2013.00303. URL: <https://www.frontiersin.org/journals/human-neuroscience/articles/10.3389/fnhum.2013.00303/full> (visited on 02/19/2026).
- Damasio, A R and D Tranel (June 1993). “Nouns and verbs are retrieved with differently distributed neural systems.” In: *Proceedings of the National Academy of Sciences* 90.11, pp. 4957–4960. DOI: 10.1073/pnas.90.11.4957. URL: <https://www.pnas.org/doi/abs/10.1073/pnas.90.11.4957> (visited on 02/26/2026).
- Damasio, H. et al. (June 2001). “Neural correlates of naming actions and of naming spatial relations”. In: *NeuroImage* 13.6, pp. 1053–1064. ISSN: 1053-8119. DOI: 10.1006/nimg.2001.0775.
- De Renzi, E. and G. di Pellegrino (Dec. 1995). “Sparing of verbs and preserved, but ineffectual reading in a patient with impaired word production”. In: *Cortex; a Journal Devoted to the Study of the Nervous System and Behavior* 31.4, pp. 619–636. ISSN: 0010-9452. DOI: 10.1016/s0010-9452(13)80016-0.
- Deerwester, Scott et al. (1990). “Indexing by latent semantic analysis”. In: *Journal of the American Society for Information Science* 41.6, pp. 391–407. DOI: [https://doi.org/10.1002/\(SICI\)1097-4571\(199009\)41:6<391::AID-ASI1>3.0.CO;2-9](https://doi.org/10.1002/(SICI)1097-4571(199009)41:6<391::AID-ASI1>3.0.CO;2-9). eprint: <https://asistdl.onlinelibrary.wiley.com/doi/pdf/10.1002/%28SICI%291097-4571%28199009%2941%3A6%3C391%3A%3AAID-ASI1%3E3.O.CO%3B2-9>. URL: <https://asistdl.onlinelibrary.wiley.com/doi/abs/10.1002/%28SICI%291097-4571%28199009%2941%3A6%3C391%3A%3AAID-ASI1%3E3.O.CO%3B2-9>.
- Deniz, Fatma et al. (Sept. 25, 2019). “The Representation of Semantic Information Across Human Cerebral Cortex During Listening Versus Reading Is Invariant to Stimulus Modality”. In: *Journal of Neuroscience* 39.39, pp. 7722–7736. ISSN: 0270-6474, 1529-2401. DOI: 10.1523/JNEUROSCI.0675-19.2019. URL: <https://www.jneurosci.org/content/39/39/7722> (visited on 03/07/2026).
- Devlin, Jacob et al. (May 24, 2019). *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*. DOI: 10.48550/arXiv.1810.

04805. arXiv: 1810.04805[cs]. URL: <http://arxiv.org/abs/1810.04805> (visited on 03/07/2026).
- Fedorenko, Elena et al. (2010). “New method for fMRI investigations of language: defining ROIs functionally in individual subjects”. In: *Journal of Neurophysiology* 104.2.
- Fischl, Bruce et al. (1999). “High-resolution intersubject averaging and a coordinate system for the cortical surface”. In: *Human Brain Mapping* 8.4, pp. 272–284. DOI: [https://doi.org/10.1002/\(SICI\)1097-0193\(1999\)8:4<272::AID-HBM10>3.0.CO;2-4](https://doi.org/10.1002/(SICI)1097-0193(1999)8:4<272::AID-HBM10>3.0.CO;2-4). eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1002/%28SICI%291097-0193%281999%298%3A4%3C272%3A%3AAID-HBM10%3E3.0.CO%3B2-4>. URL: <https://onlinelibrary.wiley.com/doi/abs/10.1002/%28SICI%291097-0193%281999%298%3A4%3C272%3A%3AAID-HBM10%3E3.0.CO%3B2-4>.
- Huth, Alexander G. et al. (Apr. 28, 2016). “Natural speech reveals the semantic maps that tile human cerebral cortex”. In: *Nature* 532.7600, pp. 453–458. ISSN: 0028-0836. DOI: 10.1038/nature17637. URL: <https://pmc.ncbi.nlm.nih.gov/articles/PMC4852309/> (visited on 03/07/2026).
- Jain, Shailee and Alexander Huth (2018). “Incorporating Context into Language Encoding Models for fMRI”. In: *Advances in Neural Information Processing Systems*. Vol. 31. Curran Associates, Inc. URL: https://papers.nips.cc/paper_files/paper/2018/hash/f471223d1a1614b58a7dc45c9d01df19-Abstract.html (visited on 03/07/2026).
- Lund, Kevin and Curt Burgess (June 1, 1996). “Producing high-dimensional semantic spaces from lexical co-occurrence”. In: *Behavior Research Methods, Instruments, & Computers* 28.2, pp. 203–208. ISSN: 1532-5970. DOI: 10.3758/BF03204766. URL: <https://doi.org/10.3758/BF03204766>.
- Mitchell, Tom M. et al. (2008). “Predicting Human Brain Activity Associated with the Meanings of Nouns”. In: *Science* 320.5880, pp. 1191–1195. DOI: 10.1126/science.1152876. eprint: <https://www.science.org/doi/pdf/10.1126/science.1152876>. URL: <https://www.science.org/doi/abs/10.1126/science.1152876>.
- Negi, Anuja et al. (2025). “Optimizing Language Model Embeddings to Voxel Activity Improves Brain Activity Predictions”. In: *bioRxiv*. DOI: 10.1101/2025.09.18.676935. eprint: <https://www.biorxiv.org/content/early/2025/11/13/2025.09.18.676935.full.pdf>. URL: <https://www.biorxiv.org/content/early/2025/11/13/2025.09.18.676935>.
- Nunez-Elizalde, A. O., A. G. Huth, and J. L. Gallant (2019). “Voxelwise encoding models with non-spherical multivariate normal priors”. In: *Neuroimage* 482-492.197.
- Peters, Matthew E. et al. (Mar. 22, 2018). *Deep contextualized word representations*. DOI: 10.48550/arXiv.1802.05365. arXiv: 1802.05365[cs]. URL: <http://arxiv.org/abs/1802.05365> (visited on 03/06/2026).
- Schrimpf, Martin et al. (Nov. 9, 2021). “The neural architecture of language: Integrative modeling converges on predictive processing”. In: *Proceedings of the National Academy of Sciences* 118.45, e2105646118. DOI: 10.1073/pnas.2105646118. URL: <https://www.pnas.org/doi/10.1073/pnas.2105646118> (visited on 03/07/2026).
- Shapiro, Kevin and Alfonso Caramazza (May 1, 2003). In: *Trends in Cognitive Sciences* 7.5, pp. 201–206. ISSN: 1364-6613, 1879-307X. DOI: 10.1016/S1364-

- 6613(03)00060-3. URL: [https://www.cell.com/trends/cognitive-sciences/abstract/S1364-6613\(03\)00060-3](https://www.cell.com/trends/cognitive-sciences/abstract/S1364-6613(03)00060-3) (visited on 02/19/2026).
- Shapiro, Kevin A. et al. (Feb. 15, 2005). "Dissociating neural correlates for nouns and verbs". In: *NeuroImage* 24.4, pp. 1058–1067. ISSN: 1053-8119. DOI: 10.1016/j.neuroimage.2004.10.015. URL: <https://www.sciencedirect.com/science/article/pii/S1053811904006226> (visited on 02/26/2026).
- Siri, Simona et al. (May 2007). "The Neural Substrate of Naming Events: Effects of Processing Demands but not of Grammatical Class". In: *Cerebral Cortex* 18.1, pp. 171–177. ISSN: 1047-3211. DOI: 10.1093/cercor/bhm043. eprint: <https://academic.oup.com/cercor/article-pdf/18/1/171/915307/bhm043.pdf>. URL: <https://doi.org/10.1093/cercor/bhm043>.
- Tuckute, Greta et al. (Mar. 2024). "Driving and suppressing the human language network using large language models". In: *Nature Human Behaviour* 8.3, pp. 544–561. ISSN: 2397-3374. DOI: 10.1038/s41562-023-01783-7.
- Tyler, L. K et al. (Jan. 1, 2004). "Neural processing of nouns and verbs: the role of inflectional morphology". In: *Neuropsychologia* 42.4, pp. 512–523. ISSN: 0028-3932. DOI: 10.1016/j.neuropsychologia.2003.10.001. URL: <https://www.sciencedirect.com/science/article/pii/S0028393203002410> (visited on 02/19/2026).
- Vaswani, Ashish et al. (Aug. 2, 2023). *Attention Is All You Need*. DOI: 10.48550/arXiv.1706.03762. arXiv: 1706.03762[cs]. URL: <http://arxiv.org/abs/1706.03762> (visited on 03/07/2026).
- Vigliocco, Gabriella et al. (Jan. 1, 2011). "Nouns and verbs in the brain: A review of behavioural, electrophysiological, neuropsychological and imaging studies". In: *Neuroscience & Biobehavioral Reviews* 35.3, pp. 407–426. ISSN: 0149-7634. DOI: 10.1016/j.neubiorev.2010.04.007. URL: <https://www.sciencedirect.com/science/article/pii/S0149763410000916> (visited on 02/27/2026).
- Yokoyama, Satoru et al. (Aug. 2006). "Cortical Mechanisms Involved in the Processing of Verbs: An fMRI Study". In: *Journal of Cognitive Neuroscience* 18.8, pp. 1304–1313. ISSN: 0898-929X. DOI: 10.1162/jocn.2006.18.8.1304. eprint: <https://direct.mit.edu/jocn/article-pdf/18/8/1304/1935948/jocn.2006.18.8.1304.pdf>. URL: <https://doi.org/10.1162/jocn.2006.18.8.1304>.




Figure 5. Noun and verb sensitivity across subjects in the reading modality. See Fig 3 for details.

Supplementary Material